

Lab 3

Thursday, 22 August 2024 11:36 AM

3.1 Task 1: Understanding the process tree (20%)

Try tracing the process tree for a few other processes other than bash, in the `ptree.txt` file you created in Section 2.2. What can you observe?

Now, try running the `htop` command as follows:

```
1 $ htop -t
```

Is the result of running `htop` in agreement with your results from manually tracing the process tree earlier?

Submit the `ptree.txt` file that you have created earlier, together with a brief summary of your findings.

Running the `htop -t` command displays a lot more processes compared to the `ptree.txt` file, this is because it is displaying all the child processes too whereas `ps` only displays a selection of the active processes, `htop` also continuously updates what processes are being run.

Run the program, and look at the output and trace back the sequence of parent process IDs for the child process. **Document your findings for the following:**

- (a) Locate the child process in the listing and underline its PID;
- (b) Put a circle around the parent's process ID for that child (it's right next to the PID);
- (c) Draw a line from the circled parent's ID to the PID for that process (find it on a previous line);
- (d) Repeat this for the parent, and its parent, etc., until you reach the process with PID 1.

In addition to your findings above, are there any differences between using `exec()` and `system()`? If there are differences, can you list two of them?

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
4	S	0	<u>1</u>	0	1	80	0	-	41464	-	?	00:00:01	systemd
5	S	0	<u>2</u>	<u>1</u>	0	80	0	-	619	-	?	00:00:00	init-systemd(Ub
0	S	0	6	2	0	80	0	-	619	-	?	00:00:00	init
4	S	0	38	1	0	79	-1	-	9885	-	?	00:00:00	systemd-journal
4	S	0	58	1	0	80	0	-	5526	-	?	00:00:00	systemd-udev
1	S	0	68	1	0	80	0	-	38248	-	?	00:00:00	snapfuse
1	S	0	74	1	0	80	0	-	38248	-	?	00:00:00	snapfuse
1	S	0	76	1	0	80	0	-	94321	-	?	00:00:00	snapfuse
1	S	0	83	1	0	80	0	-	38281	-	?	00:00:00	snapfuse
1	S	0	88	1	0	80	0	-	38248	-	?	00:00:00	snapfuse
1	S	0	95	1	0	80	0	-	56972	-	?	00:00:00	snapfuse
1	S	0	100	1	1	80	0	-	75630	-	?	00:00:01	snapfuse
1	S	0	104	1	0	80	0	-	38248	-	?	00:00:00	snapfuse
4	S	101	128	1	0	80	0	-	6385	-	?	00:00:00	systemd-resolve
4	S	0	167	1	0	80	0	-	1077	-	?	00:00:00	cron
4	S	102	169	1	0	80	0	-	2146	-	?	00:00:00	dbus-daemon
4	S	0	173	1	0	80	0	-	7525	-	?	00:00:00	networkd-dispat

4 S	0	167	1	0	80	0	-	1077	-	?	00:00:00	systemd-resolve
4 S	102	169	1	0	80	0	-	2146	-	?	00:00:00	cron
4 S	0	173	1	0	80	0	-	7525	-	?	00:00:00	dbus-daemon
4 S	104	175	1	0	80	0	-	55601	-	?	00:00:00	networkd-dispat
4 S	0	180	1	0	80	0	-	514833	-	?	00:00:00	rsyslogd
4 S	0	198	1	0	80	0	-	3833	-	?	00:00:00	snapt
4 S	0	260	1	0	80	0	-	1196	-	?	00:00:00	systemd-logind
4 S	0	269	1	0	80	0	-	26806	-	?	00:00:00	subiquity-serve
4 S	0	285	1	0	80	0	-	810	-	hvc0	00:00:00	unattended-upgr
4 S	0	290	1	0	80	0	-	799	-	tty1	00:00:00	agetty
4 S	0	318	260	3	80	0	-	38530	-	?	00:00:03	agetty
4 S	0	357	2	0	80	0	-	1882	-	pts/1	00:00:00	python3.10
4 S	1000	383	1	0	80	0	-	4233	-	?	00:00:00	login
5 S	1000	384	383	0	80	0	-	42228	-	?	00:00:00	systemd
4 S	1000	389	357	0	80	0	-	1532	core_s	pts/1	00:00:00	(sd-pam)
5 S	0	400	2	0	80	0	-	621	-	?	00:00:00	bash
5 S	0	401	400	0	80	0	-	625	-	?	00:00:00	SessionLeader
4 S	1000	402	401	0	80	0	-	723	do_wai	pts/0	00:00:00	Relay(402)
0 S	1000	403	402	0	80	0	-	723	do_wai	pts/0	00:00:00	sh
0 S	1000	409	403	0	80	0	-	723	do_wai	pts/0	00:00:00	sh
0 S	1000	413	409	26	80	0	-	335388	-	pts/0	00:00:21	sh
0 S	0	590	318	1	80	0	-	11059	-	?	00:00:01	node
5 S	0	766	2	0	80	0	-	621	-	?	00:00:00	python3
5 S	0	767	766	0	80	0	-	625	-	?	00:00:00	SessionLeader
4 S	1000	768	767	0	80	0	-	249510	-	pts/2	00:00:00	Relay(768)
5 S	0	776	2	0	80	0	-	621	-	?	00:00:00	node
5 S	0	777	776	0	80	0	-	625	-	?	00:00:00	SessionLeader
4 S	1000	778	777	0	80	0	-	247859	-	pts/3	00:00:00	Relay(778)
0 S	1000	787	413	13	80	0	-	5607461	-	pts/0	00:00:07	node
0 S	1000	823	413	10	80	0	-	310174	-	pts/0	00:00:05	node

0 S	1000	845	787	19	80	0	-	38438	-	pts/0	00:00:10	node
5 S	0	847	2	0	80	0	-	621	-	?	00:00:00	cpptools
5 S	0	848	847	0	80	0	-	625	-	?	00:00:00	SessionLeader
4 S	1000	850	848	0	80	0	-	723	do_wai	pts/4	00:00:00	Relay(850)
0 S	1000	851	850	0	80	0	-	723	do_wai	pts/4	00:00:00	sh
0 S	1000	861	851	0	80	0	-	247864	-	pts/4	00:00:00	sh
0 S	1000	918	861	0	80	0	-	724	-	pts/4	00:00:00	node
0 R	1000	938	413	1	80	0	-	281044	-	pts/0	00:00:00	sh
0 S	1000	950	938	0	80	0	-	1563	core_s	pts/5	00:00:00	node
0 S	1000	1076	938	0	80	0	-	1557	do_wai	pts/6	00:00:00	bash
0 S	1000	1179	401	0	80	0	-	1064146	futex_	pts/0	00:00:00	bash
0 S	1000	1432	401	0	80	0	-	1064166	futex_	pts/0	00:00:00	cpptools-srv
0 R	1000	1587	1076	0	80	0	-	662	-	pts/6	00:00:00	cpptools-srv
1 S	1000	1588	1587	0	80	0	-	662	do_wai	pts/6	00:00:00	fork2
0 S	1000	1589	1588	0	80	0	-	724	do_wai	pts/6	00:00:00	fork2
0 R	1000	1590	1589	0	80	0	-	1872	-	pts/6	00:00:00	sh

system() will execute the supplied command in a child process that it spawns. exec() will replace the current process with the invocation of the new executable that you specify

3.3 Task 3: Understanding process scheduling (30%)

Refer to the program (fork.c) which you have downloaded in the pre-class preparation, Section 2.3. Notice the for-loop with the loop variable 'j'. This is a *delay* loop, to make the parent and child processes take a longer time to complete their work. If you haven't already done so, try running the fork program several times. **What can you say about the output?**

The order of the child and parent changes in the latter iterations is different each run of the program

Now, to see the effect of the delay loop, change the limit value by dropping one zero (to make it 100,000). Re-compile and run the program. **What has happened to the pattern of executions of the two processes?** Continue to drop one zero from the limit value until the value of the limit is 100. **Are there any changes?**

The lower the limit the more child processes are at the end, when the limit is high it alternates between parent and child processes but as the limit is lower the parent processes will finish earlier and the child processes will be left at the end

Now, what if we added more zeroes to initial limit value of 1,000,000 instead, so that the value of the limit is 10,000,000, or 100,000,000? **Does this result in anything different?**

It takes a lot longer for the program to finish running

By the way, there is an alternative (less precise but simpler) method for waiting — but in integer chunks of *seconds*. Now, change the waiting code in the program with the following bit of code:

```
1 sleep(2);
```

Re-compile and run the program again. **What differences in the output do you notice?**

Yes a pair of child and parent processes run and finish and then print rather than all of the processes printing at the same time